

CSGE602055 Operating Systems

CSF2600505 Sistem Operasi

Week 09: File System & Persistent Storage

Rahmat M. Samik-Ibrahim

University of Indonesia

<http://rms46.vlsm.org/2/207.html>

Always check for the latest revision!

REV138 14-May-2018

Operating Systems 2018-1 (Room 3114 Tue/Thu)

Class: A (10:00-12:00) | B (13:00-15:00) | C (16:00-18:00)

Week	Schedule	Topic	OSC9
Week 00	06 Feb - 12 Feb 2018	Overview 1	Ch. 1, 16
Week 01	13 Feb - 19 Feb 2018	Overview 2 & Scripting	Ch. 1, 2
Week 02	20 Feb - 26 Feb 2018	Protection, Security, Privacy, & C-language	Ch. 14, 15
Week 03	27 Feb - 05 Mar 2018	I/O, BIOS, Loader, & Systemd	Ch. 13
Week 04	06 Mar - 12 Mar 2018	Addressing, Shared Lib, & Pointer	Ch. 8
Week 05	13 Mar - 19 Mar 2018	Virtual Memory	Ch. 9
Reserved	20 Mar - 24 Mar 2018		
Mid-Term	03 Apr 2018	13:00 - 15:30 (UTS)	
Week 06	05 Apr - 11 Apr 2018	Concurrency: Processes & Threads	Ch. 3, 4
Week 07	12 Apr - 18 Apr 2018	Synchronization	Ch. 5, 7
Week 08	19 Apr - 25 Apr 2018	Scheduling	Ch. 6
Week 09	26 Apr - 07 May 2018	File System & Persistent Storage	Ch. 10, 11, 12
Reserved	08 May - 14 May 2018		
Week 10	15 May - 21 May 2018	Network Sockets Programming & I/O Programming	
Reserved	22 May - 22 May 2018		
Final	31 May 2018	13:00 - 15:00 (UAS)	
Deadline	07 Jun 2018 16:00	Extra assignment deadline	

• The Check List (Operating Systems)

- ☐ **Starting Point:** <http://rms46.vlsm.org/2/207.html>
- ☐ **Text Book:** any recent/decent OS book but map it to **OSC9**.
- ☐ Create **public** project "os181" on your github.com account.
 - ☐ Create file "README.md" and add an extra line every week. For e.g.¹:
ZCZC Sistem Operasi 2018 Awal (1)
ZCZC W01 Have tried demo for week 01.
ZCZC W02 Week 02 is done.
ZCZC W03 Week 03 is done.
- ☐ Encode your **QRC** with image size of approximately 250x250 pixels:
"OS181 CLASS ID GITHUB-ACCOUNT SSO-ACCOUNT SIAK-Full-Name"
Special for Week 00: Mail your **embedded** QRC to: os181@vlsm.org
with Subject: [W00] CLASS ID SIAK-NAME.
- ☐ Write your Memo (with QRC) **every week**.
- ☐ Using your **SSO** account, login to badak.cs.ui.ac.id via kawung.cs.ui.ac.id.
 - ☐ Check folder badak:///extra/Week00/
 - ☐ Every week, copy the weekly demo files to your own home directory.
Eg. for Week00:
cp -r /extra/Week00/W00-demos/ W00-demos/

¹Week 00 line is optional. The following "ZCZC WXX" weekly tags are mandatory.

Agenda

- 1 Start
- 2 Agenda
- 3 Mass-Storage Structure
- 4 Disk Management
- 5 RAID
- 6 File System Interface
- 7 File System Implementation
- 8 Devices
- 9 FUSE
- 10 The End

Week 09: Persistent Storage & File System

- Reference: (OSC9-ch10 OSC9-ch11 OSC9-ch12 demo-w09)
- Mass-Storage Structure
 - Obsolete: Magnetic Tape, Disket
 - Until When?: Magnetic Disk, DVD
 - Until When?: Mechanical Disk Arm Scheduling
 - Solid-State Disks (SSD)
 - (What is a) Flash Disk
- Attached-Storage
 - Host-Attached Storage: via I/O
 - Network-Attached Storage (NAS): via distributed FS
 - Storage Area Network (SAN): dedicated Network
- Legacy Linux I/O Scheduling Algorithm.
 - Deadline Scheduler
 - Completely Fair Queueing (CFQ)

Disk Management

- Formatting
 - Low Level (Physical)
 - High Level (FS)
- Boot Block
- Disk Partition
 - "MBR"-scheme
 - upto 4 primary partition
 - upto 2 TB disk
 - "GPT"-scheme
 - "unlimited" partition
 - "unlimited" disk
 - redundancy
- Swap Space Management: On Partition or FS?

RAID: Redundant Array of In* Disks

- RAID 0, 1, 5, 6, 10, 100
- Note (<http://www.commodore.ca/windows/raid5/raid5.htm>):
 - RAID was created to enhance data performance, reliability and availability.
 - Striping, parity checking and mirroring are three primary functions of RAID systems.
 - RAID performs its functions transparent to the operating system.
 - Systems are typically defined by ranks consisting of five disks each connected to one or two Disk Array Controllers.
 - Different RAID levels provide varying degrees of speed and data protection.
- Problems with RAID
- Stable-Storage Implementation

File System Interface

- File Concept

- File Attributes: Name, Id, Type, Location, Size, Protection, Time Stamp: create, last modified, last accessed.
- File Operation
 - Create/Delete/Truncate
 - Open/Close
 - Read/Write
- File Types: Executable, Object, Source Code, Library, Markup, Markdown, Archive, Compressed.
- File Structure
- Access Methods: Sequential vs Direct Access

- Directory and Disk Structure

- Three-Structured Directories
- FS Mounting vs. Volume Based System

- File Sharing

- Protection: Access Control

File Systems Implementation

- File System Layers / Structure
 - Application Programs
 - Logical File Systems
 - File-Organization Module
 - Basic File Systems
 - I/O Control
 - Hardware Device
- File System Implementation
- File Control Block
- FS In Memory Structure
- VFS: Virtual File Systems
 - How to support multiple File Systems
 - I.e. How to support multiple `open()/close()` `read()/write()` operations

Implementation and Allocation Method

- Directory Implementation
 - Linear List
 - Hash Table
- Allocation Method
 - Contiguous
 - Linked
 - Indexed
 - Combined Scheme
- Free Space Management
- Efficiency & Performance
- Recovery

- the /dev/ directory
 - /etc/fstab: configuration of filesystems
 - /etc/mtab → /proc/mounts: mounted filesystems
 - /proc/swaps: swap filesystems
 - df: checking disk space and filesystems
 - Device Major and Minor Numbers
 - UUID - Universally Unique Identifier (128 bits)
 - GUID - Globally Unique Identifiers: `ls -al /dev/disk/by-uuid`
 - practically is NOT guaranteed unique
 - FUSE: Filesystem in Userspace
 - BBFS: Big Brother File System
- More Storage Structure
 - tmpfs
 - objfs
 - ctfs
 - lofs
 - procfs
 - ufs
 - zfs

FUSE

```
>>>>> $ ls -al
total 172
drwxr-xr-x 4 demo demo 4096 Apr 25 17:30 .
drwx----- 4 demo demo 4096 Apr 25 17:30 ..
-rw-r--r-- 1 demo demo 2214 Apr 25 17:30 1-READ-THIS-FIRST.txt
drwxr-xr-x 2 demo demo 4096 Apr 25 17:30 disk-images
drwxr-xr-x 5 demo demo 4096 Apr 25 17:30 fuse-tutorial-2018-02-04
-rw-r--r-- 1 demo demo 151814 Apr 25 17:30 fuse-tutorial-2018-02-04.tgz
```

```
>>>>> $ cat 1-READ-THIS-FIRST.txt
```

This demo is available in badak:///extra/.

See also: <https://github.com/UI-FASILKOM-OS/os181>.

You should copy this folder to your own working folder.

DO NOT work inside the /extra folder!

=====

Folder Name:

Week09/

FUSE (2)

To copy the folder to your home directory:

```
cp -r /extra/Week09/W09-demos/ W09-demos/
```

=====

File Listing:

- * 1-READ-THIS-FIRST.txt (this file)
- * disk-images/ (directory)
- * fuse-tutorial-2018-02-04 (fuse tutorial)
- * fuse-tutorial-2018-02-04.tgz (source).

A. disk-images

```
cd disk-images/
```

=====

B. fuse-tutorial

1. UBUNTU's deb packages (privilege):

```
sudo apt-get install autoconf automake build-essential \  
fuse libfuse-dev pkg-config sshfs
```

2. Get the tarball with

```
wget http://www.cs.nmsu.edu/~pfeiffer/fuse-tutorial.tgz
```

3. List and open the tarball with

```
tar tfz fuse-tutorial.tgz  
tar xzf fuse-tutorial.tgz
```

4. Enter the directory (yours may be a different version)

```
cd fuse-tutorial-2018-02-04/  
ls -al
```

FUSE (3)

5. Read the manual with
lynx index.html

Writing a FUSE Filesystem: a Tutorial

Joseph J. Pfeiffer, Jr., Ph.D. (pfeiffer@cs.nmsu.edu)
Emeritus Professor
Department of Computer Science, New Mexico State University

Version of 2018-02-04

One of the real contributions of Unix has been the view that "everything is a file". A tremendous number of radically different sorts of objects, from data storage to file format conversions to internal operating system data structures, have been mapped to the file abstraction.

One of the more recent directions this view has taken has been Filesystems in User Space, or FUSE (no, the acronym really doesn't work. Oh well). The idea here is that if you can envision your interaction with an object in terms of a directory structure and filesystem operations, you can write a FUSE file system to provide that interaction. You just write code that implements file operations like `open()`, `read()`, and `write()`; when your filesystem is mounted, programs are able to access the data using the standard file operation system calls, which call your code.

FUSE filesystems have been written to do everything from providing remote access to files on a different host without using NFS or CIFS (see SSHFS at [2]<https://github.com/libfuse/sshfs>) to implementing a filesystem to talk to devices using the Media Transfer protocol (see

[.....])

FUSE (4)

```
6. Run
   ./configure
   make
```

```
7 cd example
```

TO TRY:

```
$ ls -al rootdir
$ ls -al mountdir
$ df
$ ../src/bbfs rootdir/ mountdir/
$ df
$ ls -al rootdir
$ ls -al mountdir
```

TO PLAY:

```
$ cd mountdir
$ touch blah-blah-blah.txt
$ ls -al
$ cd ..
$ ls -al rootdir
```

TO FINISH:

```
$ fusermount -u mountdir
```

EXTRA:

```
# /etc/fstab: configuration of filesystems
# /etc/mtab --> /proc/mounts: mounted filesystems
# /proc/swaps: swap filesystems
# df: checking disk space and filesystems
# GUID (Globally Unique Identifiers) ls -al /dev/disk/by-uuid
RMS
```

FUSE (5)

```
>>>> $ ./configure
checking for a BSD-compatible install... /usr/bin/install -c
checking whether build environment is sane... yes
checking for a thread-safe mkdir -p... /bin/mkdir -p
checking for gawk... gawk
checking whether make sets $(MAKE)... yes
checking whether make supports nested variables... yes
checking for gcc... gcc
checking whether the C compiler works... yes
checking for C compiler default output file name... a.out
checking for suffix of executables...
checking whether we are cross compiling... no
checking for suffix of object files... o
checking whether we are using the GNU C compiler... yes
checking whether gcc accepts -g... yes
checking for gcc option to accept ISO C89... none needed
checking whether gcc understands -c and -o together... yes
checking for style of include used by make... GNU
checking dependency style of gcc... gcc3
checking how to run the C preprocessor... gcc -E
checking for grep that handles long lines and -e... /bin/grep
checking for egrep... /bin/grep -E
checking for ANSI C header files... yes
[...]
checking for fdatasync... yes
checking that generated files are newer than configure... done
configure: creating ./config.status
config.status: creating Makefile
config.status: creating html/Makefile
config.status: creating src/Makefile
config.status: creating src/config.h
config.status: executing depfiles commands
```


FUSE (6)

```
>>>>> $ make
Making all in example
make[1]: Entering directory '/home/demo/mydemo/W09-demos/fuse-tutorial-2018-02-04/example'
mkdir -p mountdir
mkdir -p rootdir
echo "bogus file" > rootdir/bogus.txt
make[1]: Leaving directory '/home/demo/mydemo/W09-demos/fuse-tutorial-2018-02-04/example'
Making all in html
make[1]: Entering directory '/home/demo/mydemo/W09-demos/fuse-tutorial-2018-02-04/html'
make[1]: Nothing to be done for 'all'.
make[1]: Leaving directory '/home/demo/mydemo/W09-demos/fuse-tutorial-2018-02-04/html'
Making all in src
make[1]: Entering directory '/home/demo/mydemo/W09-demos/fuse-tutorial-2018-02-04/src'
make all-am
make[2]: Entering directory '/home/demo/mydemo/W09-demos/fuse-tutorial-2018-02-04/src'
gcc -DHAVE_CONFIG_H -I. -D_FILE_OFFSET_BITS=64 -I/usr/include/fuse -g -O2 -MT bbfs.o -MD -MP -MF
.deps/bbfs.Tpo -c -o bbfs.o bbfs.c
mv -f .deps/bbfs.Tpo .deps/bbfs.Po
gcc -DHAVE_CONFIG_H -I. -D_FILE_OFFSET_BITS=64 -I/usr/include/fuse -g -O2 -MT log.o -MD -MP -MF
.deps/log.Tpo -c -o log.o log.c
mv -f .deps/log.Tpo .deps/log.Po
gcc -D_FILE_OFFSET_BITS=64 -I/usr/include/fuse -g -O2 -o bbfs bbfs.o log.o -lfuse -pthread
make[2]: Leaving directory '/home/demo/mydemo/W09-demos/fuse-tutorial-2018-02-04/src'
make[1]: Leaving directory '/home/demo/mydemo/W09-demos/fuse-tutorial-2018-02-04/src'
make[1]: Entering directory '/home/demo/mydemo/W09-demos/fuse-tutorial-2018-02-04'
make[1]: Nothing to be done for 'all-am'.
make[1]: Leaving directory '/home/demo/mydemo/W09-demos/fuse-tutorial-2018-02-04'
>>>>> $
```

FUSE (7)

```
>>>> $ cd example/
>>>> $ ls -al rootdir/
total 12
drwxr-xr-x 2 demo demo 4096 Apr 25 18:23 .
drwxr-xr-x 4 demo demo 4096 Apr 25 18:23 ..
-rw-r--r-- 1 demo demo 11 Apr 25 18:23 bogus.txt
>>>> $ ls -al mountdir/
total 8
drwxr-xr-x 2 demo demo 4096 Apr 25 18:23 .
drwxr-xr-x 4 demo demo 4096 Apr 25 18:23 ..
>>>> $ df
Filesystem      1K-blocks      Used Available Use% Mounted on
udev              10240          0       10240   0% /dev
tmpfs            1639412      103116    1536296   7% /run
/dev/vda2         9515660     1677648    7331596  19% /
/dev/vdc1        32895760    12093508    19108200  39% /usr
tmpfs             4098528          0     4098528   0% /dev/shm
tmpfs             5120          0        5120   0% /run/lock
tmpfs            4098528          0     4098528   0% /sys/fs/cgroup
/dev/vdb1        515929528   38454128   451244668   8% /home
tmpfs            819708          0      819708   0% /run/user/1002
>>>> $ ../src/bbfs rootdir/ mountdir/
Fuse library version 2.9
about to call fuse_main
>>>> $ df
Filesystem      1K-blocks      Used Available Use% Mounted on
udev              10240          0       10240   0% /dev
[...]
tmpfs            819708          0      819708   0% /run/user/1002
bbfs            515929528   38454136   451244660   8% /home/demo/mydemo/W09-demos/
                                     fuse-tutorial-2018-02-04/example/mountdir
>>>> $
```

FUSE (8)

```
>>>> $ ls -al rootdir/
total 12
drwxr-xr-x 2 demo demo 4096 Apr 25 18:23 .
drwxr-xr-x 4 demo demo 4096 Apr 25 18:26 ..
-rw-r--r-- 1 demo demo 11 Apr 25 18:23 bogus.txt
>>>> $ ls -al mountdir/
total 12
drwxr-xr-x 2 demo demo 4096 Apr 25 18:23 .
drwxr-xr-x 4 demo demo 4096 Apr 25 18:26 ..
-rw-r--r-- 1 demo demo 11 Apr 25 18:23 bogus.txt
>>>> $ cd mountdir/
>>>> $ touch blah-blah-blah.txt
>>>> $ ls -al
total 12
drwxr-xr-x 2 demo demo 4096 Apr 25 18:30 .
drwxr-xr-x 4 demo demo 4096 Apr 25 18:26 ..
-rw-r--r-- 1 demo demo 0 Apr 25 18:30 blah-blah-blah.txt
-rw-r--r-- 1 demo demo 11 Apr 25 18:23 bogus.txt
>>>> $ cd ..
>>>> $ ls -al rootdir/
total 12
drwxr-xr-x 2 demo demo 4096 Apr 25 18:30 .
drwxr-xr-x 4 demo demo 4096 Apr 25 18:26 ..
-rw-r--r-- 1 demo demo 0 Apr 25 18:30 blah-blah-blah.txt
-rw-r--r-- 1 demo demo 11 Apr 25 18:23 bogus.txt
>>>> $ fusermount -u mountdir
>>>> $ ls -al mountdir/
total 8
drwxr-xr-x 2 demo demo 4096 Apr 25 18:23 .
drwxr-xr-x 4 demo demo 4096 Apr 25 18:26 ..
>>>> $
```

The End

- ☐ This is the end of the presentation.
- ☒ This is the end of the presentation.
 - This is the end of the presentation.